Auditor loop: detect → localize → name → mitigate

Confound-agnostic core (region box, or (confound, label) pair) — serves synthetic tag, real tube, and cross-site from one implementation

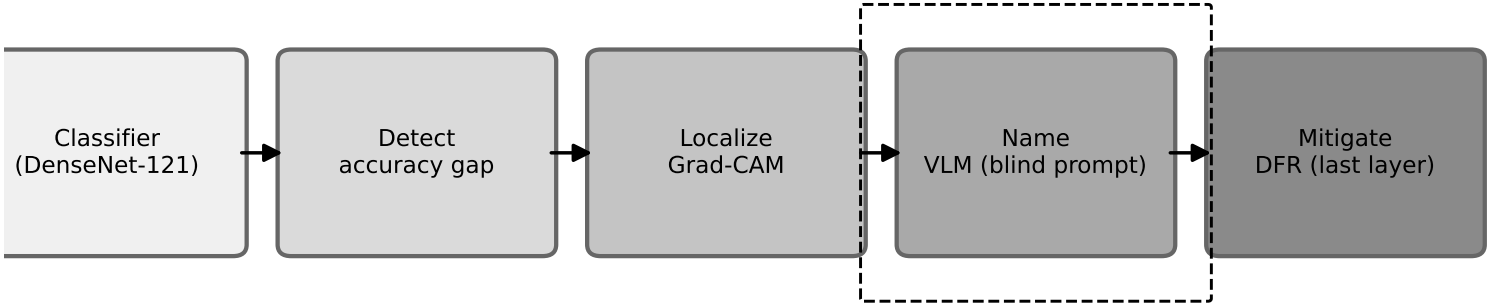

In plain terms: the classifier's decision is checked for a hidden shortcut, the image region driving that shortcut is located, a vision-language model is asked to name what it sees there with no hint of what to look for, and only then is the model retrained to stop relying on it.