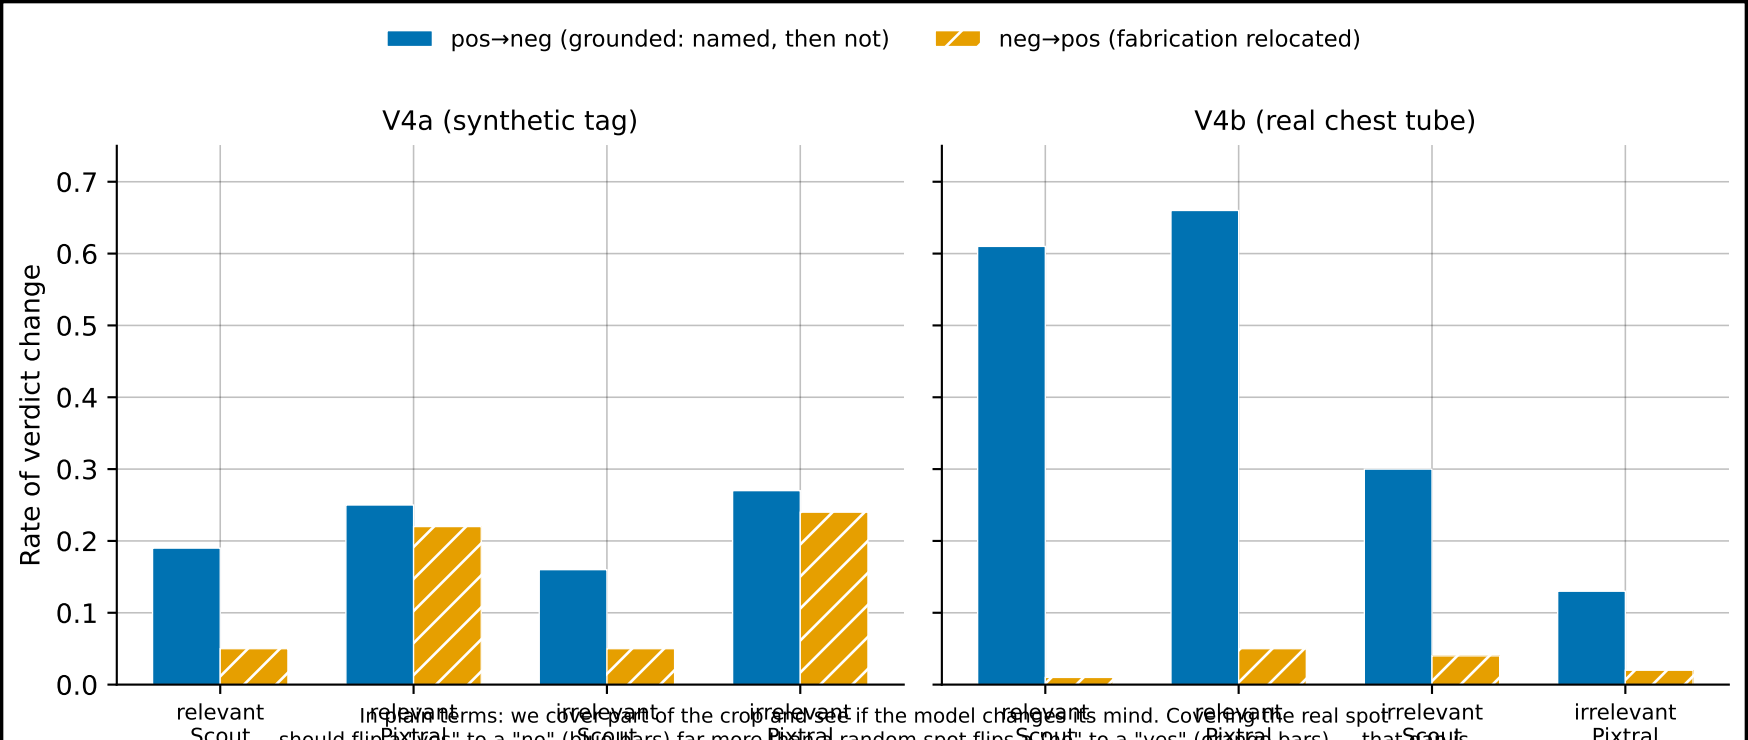

In plain terms: we cover part of the crop and see if the model changes its mind. Covering the real spot should flip a "yes" to a "no" (blue bars) far more than a random spot flips a "no" to a "yes" (orange bars) — that gap is evidence the model is actually looking at the device, not just guessing. V4a relevant-occlusion is contaminated by the gray occluder itself tripping the matcher (see §6.2 caveat); the grounding claim rests on the V4b-internal contrast (relevant → far more pos→neg than irrelevant).