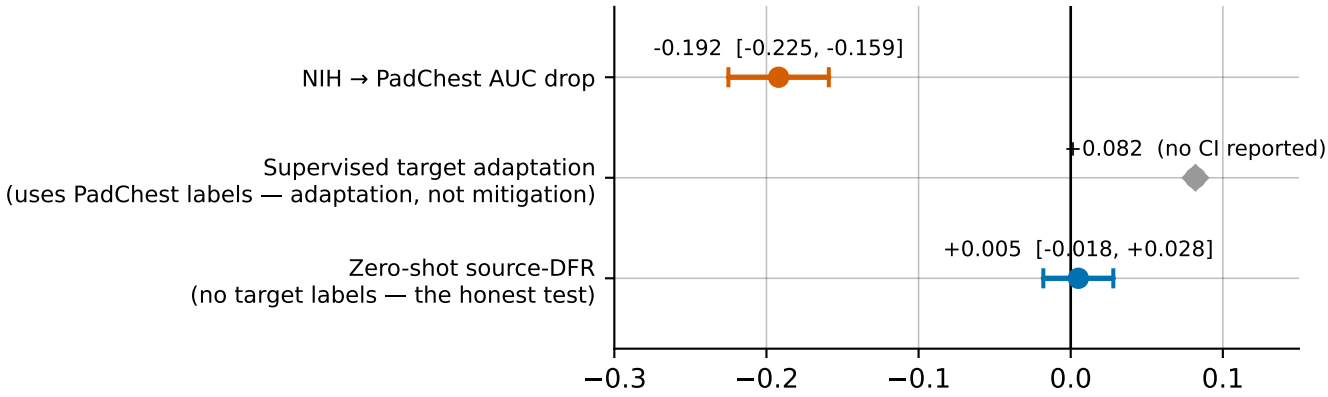

In plain terms: a classifier trained at one hospital loses about 19 points of AUC at a second hospital (top row). Letting it see the new hospital's own labels recovers most of that (middle row, dashed diamond = single-run point estimate, no CI available) — but that's cheating, since real deployments rarely have those labels. The bottom row is the honest test: a label-free fix, and its recovery is close to zero, i.e. it does not meaningfully help.